

# Exercise 1

# Exercise 1

## Recap: Local Beam Search

- Local Search
- Difference to Hill Climbing: maintain  $k$  instead of just one state
- Calculate for each maintained state the neighboring states → out of this set choose the best  $k$  states → repeat
- Problem: low diversity of the  $k$  states (all  $k$  states are very similar to each other)
- Solution: stochastic beam search
- Beam Search is often used in Natural Language Processing (NLP) (along with autoregressive models such as Transformer models)

# Exercise 1

Give the name of the algorithm resulting from the following special cases:

a) Local Beam Search with  $k = 1$

**Hill Climbing**

b) Local Beam Search with exactly one initial state and no restriction on the number of maintained states

**Breadth-first search**

# Exercise 1

Give the name of the algorithm resulting from the following special cases:

c) Simulated Annealing with always  $T = 0$

Simulated Annealing: accept next state if  $\Delta E > 0$ . Else accept the next state with probability  $e^{\Delta E/T}$ .

With  $T = 0 \rightarrow$  never accept a state if  $\Delta E \leq 0$ . This is equal to **Hill Climbing**

d) Simulated Annealing with always  $T = \infty$

With  $T = \infty \rightarrow$  Next state is always accepted. This is equal to a **Random Walk**

# Exercise 2

# Recap: Archetypal analysis

$X$  = Observations

$$\min RSS = \min ||X - \alpha Z^T||_2$$

with:

$Z$  = Archetypes

$$Z = X^T \beta$$

$\alpha$  = Assignment of  
observation  $i$  to  
archetype  $j$

$$\sum_{j=1}^K \alpha_{ij} = 1$$

$$\forall i \in [1, n] \\ \alpha_{ij} \geq 0$$

$\beta$  = Assignment of  
archetype  $j$  to  
observation  $i$

$$\sum_{i=1}^n \beta_{ji} = 1$$

$$\forall j \in [1, K] \\ \beta_{ji} \geq 0$$

# Archetypal analysis

*a) What is the constraint on minimizing the residual sum of squares with respect to the assignment of an observation to the archetypes?*

- For each observation, the sum of the assignments to the archetypes is equal to 1
- Alternative: Observations are **convex combinations** of the archetypes

$$\forall i \in [1, n]: \sum_{j=1}^K \alpha_{ij} = 1$$

# Archetypal analysis

*b) What is the constraint for minimizing the residual sum of squares in terms of assigning an archetype to observations?*

- For each archetype, the sum of the assignments to observations is equal to 1
- Alternative: Archetypes are **convex combinations** of the observations


$$\forall j \in [1, K]: \sum_{i=1}^n \beta_{ji} = 1$$



# Archetypal analysis

*c) Explain the structure of the Scree plot and how the elbow criterion works.*


A Scree plot relates values of the RSS to the number of archetypes  $K$



# Archetypal analysis

*c) Explain the structure of the Scree plot and how the elbow criterion works.*

- Elbow criterion: method to determine appropriate value for  $K$
- It is a tradeoff between "few  $K$ " and "small error"
- From a certain  $K$  the RSS does not improve significantly




# Archetypal analysis

*d) In Archetypal Analysis, the number  $K$  of archetypes must be known before execution. Give three ways how  $K$  can be determined.*

1. Elbow criterion
2. Determined by the user: E.g. the user wants 3 alternative routes
3.  $K$  is already known. E.g.: Striker, midfielder, defender


# Archetypal analysis

*e,f) Given is the following 2D point cloud:*




draw  $K$  archetypes each with  $K = 1, 3, 4, 5$  where you reasonably expect them to be (an estimate is enough). Then draw the convex hull. Which number of archetypes  $K$  seems to be best suited?

# Archetypal analysis




$K = 1$

# Archetypal analysis




$K = 1$

# Archetypal analysis



$K = 3$


# Archetypal analysis



$K = 3$




# Archetypal analysis




$K = 3$

# Archetypal analysis




$K = 4$

# Archetypal analysis




$K = 4$

# Archetypal analysis




$K = 4$

# Archetypal analysis




$K = 5$

# Archetypal analysis



$K = 5$

# Archetypal analysis



$K = 5$

# Archetypal analysis

f) Which number of archetypes  $K$  seems to be best suited for the data set at hand and why?

$K = 4$ , because of the elbow criterion



# Exercise 3

# Exercise 3

Which of the following statements are true?

- a) Simulated Annealing maintains multiple solution candidates at the same time

No. Just one.

- b) Evolutionary algorithms maintain multiple solution candidates at the same time

Yes.

# Exercise 3

c) Simulated Annealing guarantees to find the global optimum at infinite runtime

Depends on the decay of  $T$ .

d) The temperature  $T$  in SA controls the tradeoff between exploration and exploitation

Yes.

# Exercise 4

# Recap: Game Theory

- Pure Strategy
  - A player selects deterministically one of the possible actions
- Mixed Strategy
  - A probability distribution over all possible actions
  - A mixed strategy is a probability distribution one uses to randomly choose among available actions in order to avoid being predictable.
  - E.g.:  $A = \{0.4, 0.6\}$
- Dominant Strategy:
  - a strategy is dominant over another iff it is the best strategy regardless of the strategy of the opponent
- Nash equilibrium:
  - No player improves by individually changing its strategy

# Exercise 4

*Let the following game be given:*

	B1	B2
A1	A:5, B:6	A:7, B:2
A2	A:4, B:5	A:9, B:1

*Decide whether the following statements are true or false:*

*a) Player A has a dominant strategy*

No! Why not?

# Exercise 4

*Let the following game be given:*

	B1	B2
A1	A:5, B:6	A:7, B:2
A2	A:4, B:5	A:9, B:1

*Decide whether the following statements are true or false:*

*b) Player B has a dominant strategy*

Yes: Move 1; no matter if player A chooses move 1 or 2, it is always best for player B to choose move 1.

# Exercise 4

*Let the following game be given:*

	B1	B2
A1	A:5, B:6	A:7, B:2
A2	A:4, B:5	A:9, B:1

*Decide whether the following statements are true or false:*

*b) Player B has a dominant strategy*

Yes: Move 1; no matter if player A chooses move 1 or 2, it is always best for player B to choose move 1:  $6 > 2$  and  $5 > 1$



# Exercise 5

# Exercise 5

*Calculate the expected rewards for the following game with the mixed strategies:*

$$A = \{0.4, 0.6\}$$

$$B = \{0.9, 0.1\}$$

	A1	A2
B1	A:5, B:3	A:4, B:6
B2	A:0, B:4	A:2, B:2

Answer: see whiteboard

# Exercise 5

*Calculate the Nash equilibrium in mixed strategies for the following game:*

	B1	B2
A1	A:4, B:4	A:1, B:0
A2	A:2, B:2	A:6, B:9

Answer: see whiteboard

# Exercise 6

# Exercise 6

*Let the following game be given (the values represent earnings in million €):*

	Mercedes, Ads	Mercedes, no Ads
BMW, Ads	2, 4	4, 2
BMW, no Ads	0.5, 9	3, 6

a) What is the Nash equilibrium (in pure strategies) of the above game?

The Nash equilibrium is: (Ads, Ads)

# Exercise 6

*Let the following game be given (the values represent earnings in million €):*

	Mercedes, Ads	Mercedes, no Ads
BMW, Ads	2, 4	4, 2
BMW, no Ads	0.5, 9	3, 6

- b) If both companies could commit not to advertise by entering into a binding contract, should they do so? If so, why?

Yes, because more efficient for both companies ( $3 > 2$  and  $6 > 4$ )

# Exercise 6

*c) Implement a program in Python which finds a Nash Eq assuming pure strategies.*

See PyCharm

# Exercise 6

*d) Find a game with 3 players and 2 actions with no Nash Eq. Verify your solution using the program from c)*

[a1,a2,a3]	r1	r2	r3
[0,0,0]			
[0,0,1]			
[0,1,0]			
[0,1,1]			
[1,0,0]			
[1,0,1]			
[1,1,0]			
[1,1,1]			



# Exercise 6

*d) Find a game with 3 players and 2 actions with no Nash Eq. Verify your solution using the program from c)*

[a1,a2,a3]	r1	r2	r3
[0,0,0]	1	0	0
[0,0,1]	1	0	0
[0,1,0]	0	1	1
[0,1,1]	1	1	0
[1,0,0]	0	1	0
[1,0,1]	0	1	1
[1,1,0]	1	1	0
[1,1,1]	1	0	1

# Exercise 7

# Exercise 7

*The four most important components of an Evolutionary Algorithm are:*

- *Evaluation*
- *Selection*
- *Recombination*
- *Mutation*

*Consider the Maximum Clique Problem. How would you implement these four components for this problem?*

# Exercise 7

*Evaluation for Maximum Clique states?*

```
def objective_maxClique(x):  
    if x is a clique in G:  
        return sum(x)  
    else:  
        return 0
```

# Exercise 7

*Selection for Maximum Clique states?*

trivial.

Just choose the states with the highest objective function value.

# Exercise 7

*Recombination for Maximum Clique states?*

$$C1 = [1, 0, 1, 0, 0, 0]$$

$$C2 = [0, 0, 1, 1, 0, 1]$$

# Exercise 7

*Recombination for Maximum Clique states?*

$$C1 = [1, 0, 1, 0, 0, 0]$$

$$C2 = [0, 0, 1, 1, 0, 1]$$

$$C3 = [1, 0, 1, 1, 0, 1]$$


# Exercise 7

*Recombination for Maximum Clique states?*

**C1 = [1, 0, 1, 0, 0, 0]**

C2 = [0, 0, 1, 1, 0, 1]

C3 = [1, 0, 1, 1, 0, 1]



Value: 2




# Exercise 7

*Recombination for Maximum Clique states?*

$C1 = [1, 0, 1, 0, 0, 0]$

$C2 = [0, 0, 1, 1, 0, 1]$

$C3 = [1, 0, 1, 1, 0, 1]$



Value: 0


# Exercise 7

*Recombination for Maximum Clique states?*

$C1 = [1, 0, 1, 0, 0, 0]$

$C2 = [0, 0, 1, 1, 0, 1]$

$C3 = [1, 0, 1, 1, 0, 1]$



Value: 0

# Exercise 7

*Mutation for Maximum Clique states?*

Randomly flip entries.

Example (1 mutation):

$$C = [1, 0, 0, 1, 0, 0]$$

$$C_{mutated} = [0, 0, 0, 1, 0, 0]$$

# Exercise 7

*Mutation for Maximum Clique states?*

Randomly flip entries.


Example (2 mutations):

$$C = [1, 0, 0, 1, 0, 0]$$

$$C_{mutated} = [0, 1, 0, 1, 0, 0]$$

# Outlook: Advanced Multi-Agent Systems

# Multi-Agent Reinforcement Learning



# MDP for Multi-Agent

- Let  $n$  be the number of agents
- $S$  is the observation space
- $A_i, i = 1, 2, \dots, n$  is the action space for each agent
- The joint action space for all agents is defined by  $A = A_1 \times A_1 \times \dots \times A_n$
- The state transition probability function is represented by  $p : S \times A \times S \rightarrow [0, 1]$
- The reward function is specified as  $r : S \times A \times S \rightarrow \mathbb{R}$
- The value function of each agent is dependent on the joint action and joint policy, which is characterized by  $V_\pi : S \times A \rightarrow \mathbb{R}$

# Some Axes of MARL

## Centralized:

- One brain / algorithm deployed across many agents
- In Actor-Critic Architectures:  
Centralized Critic, but decentralized Actors

## Decentralized:

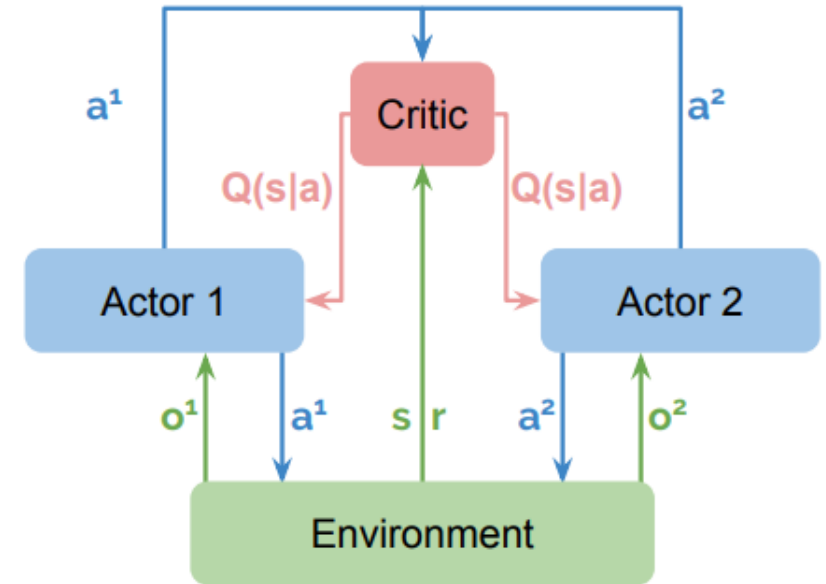
- All agents learn individually
- Treat other agents as being part of the environment → Markov property becomes invalid since the environment is no longer stationary

## Cooperative:

- Agents work together to maximize the joint reward

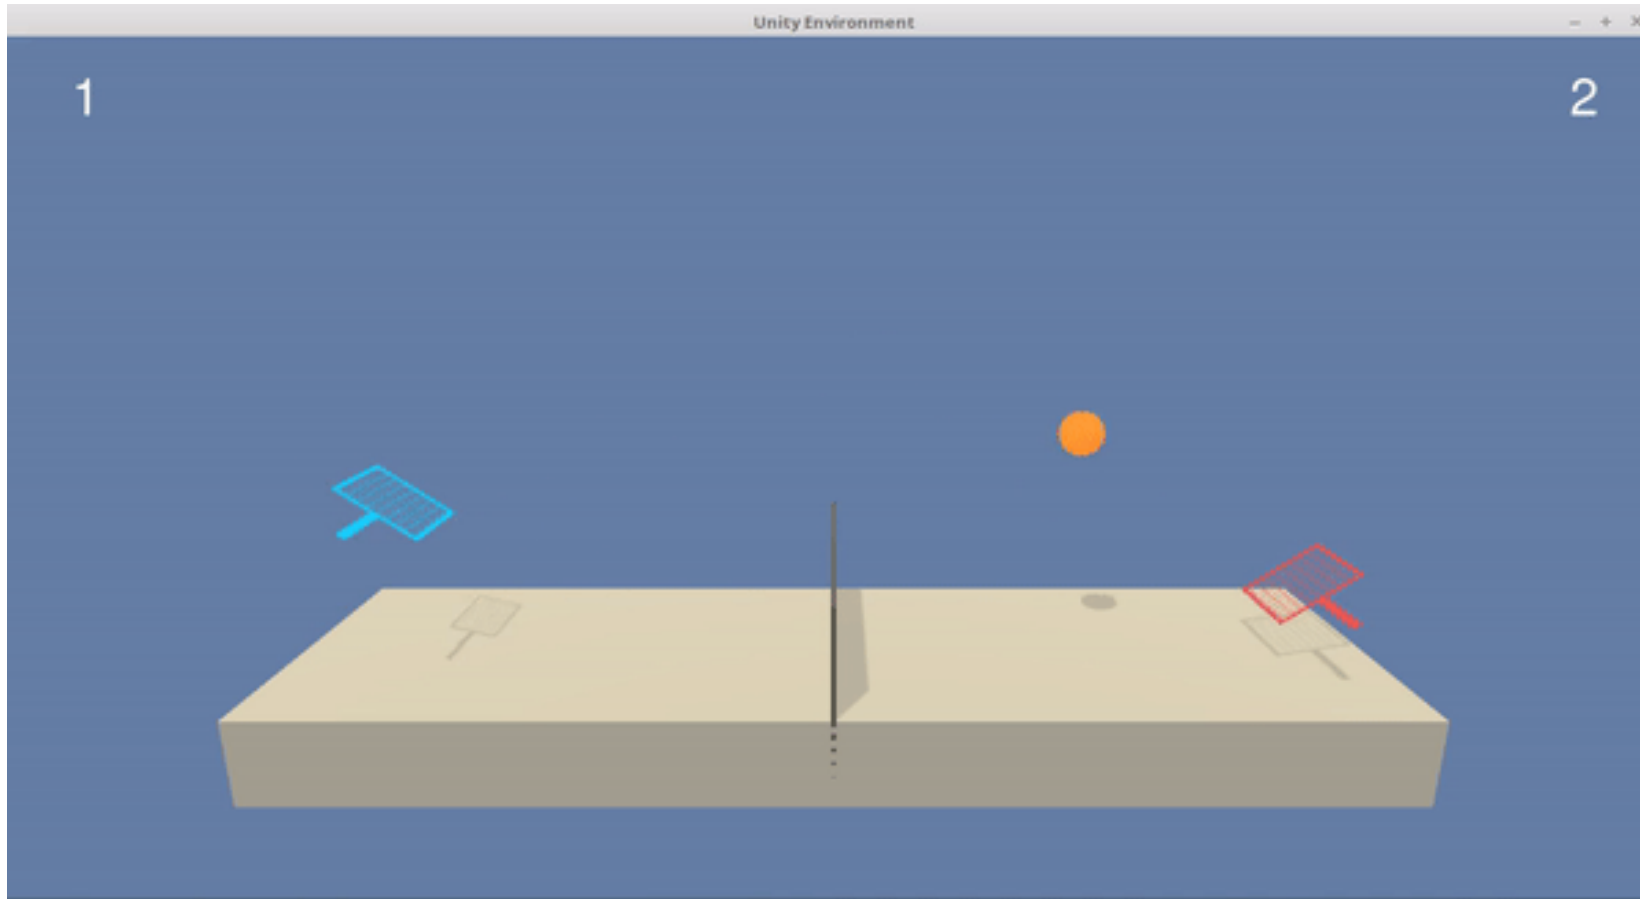
## Competitive:

- Agents trying to maximize solely their own (individual) reward






# Pure Cooperative / Competitive MARL



# Mixed-Cooperative-Competitive MARL



- Another example:  
Football
- Or any other game with  
2 teams

